

Lab 1: Prolog 1

Installing Prolog

For these labs, we will use the SWI-Prolog distribution.

macOS / OS X

Install [homebrew](#) then run `brew install swi-prolog`. You can now run the Prolog shell from the command line with `swipl`.

Important: do not install from the official website, you will get crashes.

Windows

Install from the [official website](#) (you need at least a *typical* installation). Then run the prolog shell from the start menu.

Linux

Install from your system package repositories (debian/ubuntu/mint: `swi-prolog`) or follow [the instructions](#) from the official website. Run the Prolog shell from the command line with `swipl`.

Basic Navigation

At the shell, hitting `ctrl-c` (twice for Unix) will bring up a menu with option. You can use `e` to exit.

The current directory needs to be the directory where you keep your source files. The first option is to run `swipl` there directly (Unix only). Alternatively, you can input `cd("C:/Prolog")` (relative paths also work, tab completion is available). `ls.` and `pwd.` will work like the shell `ls` and `pwd` commands.

If you need to change the current directory in your programs, use `working_directory/2` instead.

Facts

Prolog operates over a *clause database*. There are two types of clauses: *facts* and *rules*. Facts are statements which we know to be true. Here are few example facts:

```
prolog_is_simple.  
life_is_beautiful.  
father(george, maria).  
male(george).
```

Simple terms like `prolog_is_simple` or `george` are called *atoms*. The third and fourth lines exhibits complex terms (or *structures*). The facts on these lines tell us about terms or relations between terms: i.e. George is Maria's father and George is a male. In these facts, the thing before the parentheses is called the *predicate*. Predicate are often designated by their name followed their number of parameters (because overloading on the number of parameters is allowed): `father/2` and `male/1` for our exemple.

For now, facts must be defined in a file. Later we will see how we *sometimes* can define new facts using the prolog shell.

To load facts defined in a file, use the `consult/1` predicate. The course website holds such a file, named `family.pl`. Load it in one of the following ways, assuming the current directory contains the file:


```
?- consult("family.pl").  
?- consult(family).  
?- ["family.pl"].  
?- [family].
```

The `?-` sigil denotes the prolog shell. This shell can be used to load definitions from files and formulate queries over the database. As we will see later, **queries and facts are not the same thing**.

In the rest of this lab, you will have to write both definitions (facts and rules) as well as queries. We suggest that you create a file to enter all your definitions (e.g. `definitions.pl`), and that you record all your queries in a **separate** file. Load this file like the `family.pl` file above. Whenever you add a new definition to your file or change an existing definition, you can reload it by running the `make/0` predicate:

```
?- make.
```

The `family.pl` file defines a set of facts that is equivalent to the following family tree:



Exercise [facts]

Try to write some facts about our family tree in terms of `father/2` and `mother/2`:

- Maria is Alexandra's mother.
- Roger is Nadia's father.
- Laura is Anthony's mother.

Do not write these facts to your file as they would conflict with the definitions inside `family.pl`.

Logical Variables and Unification

In Prolog, objects starting with a capital letter are *logical variables*. These are not variables in the sense of most programming languages: they are not a label to some place to store data, but rather placeholders for values which are not yet known, like mathematical variables. Once their value is learnt, it will not change within a clause (excepted when backtracking). A variable can be either free, bound to a term (`x = 5.`) or aliased to another variable (`x = y.`).

Prolog uses unification, and advanced kind of pattern matching to define equality. For instance:

```
?- X = X.
```

```
true.  
  
?- 9 + 5 = 14.  
false.  
  
?- father(george, maria) = father(george, maria).  
true.  
  
?- father(george, maria) = father(george, howard).  
false.
```

No evaluation is performed on the members, they are simply defined as equals if their structures match. Prolog can also *unify* terms: two terms can be unified if it is possible to instantiate their variables in a way that both terms are equal. For example:

```
?- father(X,maria) = father(george, Y).  
X = george,  
Y = maria.
```

The terms `father(X, maria)` *unifies* with `father(george, Y)` when `X = george` and `Y = maria`.

The "result" of unification is thus a set of bindings for the variables that were unified.

Queries

Once you have fed Prolog some facts, you can start making queries. For instance:

```
?- father(george, maria).  
true.
```

Queries are not limited to yes-or-no answers. For instance, if you wish to know who is Maria's father:

```
?- father(Who, maria).  
Who = george.
```

Sometimes, queries are also called *goals* especially when they contain no

variables. *Goal* also refers to any predicate instantiation that we are trying to prove is true during the execution of a prolog program.

In Prolog a comma `,` denotes a logical *and* whereas a semi-colon denotes a logical *or*. Note that like in most programming languages, *and* binds more tightly than *or*, hence `false, true ; true.` is the same as `(false, true) ; true.`

Here is a query for all people who are male or are the children of George for instance:

```
?- male(Who) ; father(george, Who).
```

In this case there are many possible answers. You can use the space key to obtain more results and the return key to stop getting results. Prolog has multiple useful built-in predicates to handle these cases, like the `findall/3`, `bagof/3` and `setof/3` and `forall/2` predicates.

The `findall/3` predicate will unify its third parameter with a list of possible values for its first parameter, given that the second parameter is satisfied. For instance (notice the repetitions):

```
?- findall(Father, father(Father, X), List).  
List = [george, george, george, george, albert, albert, albert, roger, roger]
```

Exercise [bags]

What are the two differences between the `findall/3` and the `bagof/3` predicate? You can use the `help/1` predicate to get help about other predicates: e.g.

`help(findall/3).` or even `help(findall)`.

Use `bagof/3` to write a query that returns the same result as `findall(Father, father(Father, X), List).`

Hint: you will need some special syntax mentioned in the help.

The `setof/3` predicate works like `bagof/3` but eliminates duplicates from the result list.

The `forall/2` predicate does not provide the list of items that satisfy a query, but

rather checks whether all items matching the first query can satisfy the second query:

```
?- forall(father(albert, X), mother(maria, X)).  
true.
```

The above query will tell you whether all people who have Albert as father, also have Maria as mother.

Exercise [queries]

Write some queries to :

- know if George is the father of Tamara.
- know if Anita is the mother of Brandon.
- get all the children of Maria.
- get all the sons of Roger.
- check that Maria and Albert only have children together.
- check that all fathers are male and all mothers are female.

Rules

The second type of clauses are *rules*. Rules allow to use known facts to draw conclusions from the database. For instance:

```
male(george).           % this is a fact  
human(X) :- male(X).    % this is a rule: if X is a male, X is human
```

With that database, the query `human(george).` would return `true`.

The distinction between facts and rules is somewhat artificial: a fact

`male(george).` can be written as a rule `male(george) :- true.`. Conversely, the *turnstile* (`:-`) is a predicate and `human(X) :- male(X).` can be written `:- (human(X), male(X)).`

We also note that we can have rules without turnstiles: `unify(X, X).` is a rule defining a predicate which only succeeds if its two parameters unify. You can also see why it is important to distinguish between queries and facts, as the same

syntax can denote both a query or a rule.

Finally, we note that multiple rules may contribute to the definition of the same predicate.

So `human(X) :- man(X) ; female(X).` is equivalent to

```
human(X) :- man(X).  
human(X) :- female(X).
```

In general, you should prefer the form that uses separate rules over the `;` form.

Exercise [rules]

Write rules to describe the following relationships:

- `parent(Parent, Child)`
- `son(Son, Parent)`
- `daughter(Daughter, Parent)`
- `grandfather(Grandfather, Grandchild)`
- `grandparent(Grandparent, Grandchild)`
- `sibling(Sibling1, Sibling2)`
- `brother(Brother, Sibling)`
- `have_children_together(Person1, Person2)`
- `uncle(Uncle, Person)`

Pay attention to the details. For instance, we do not want for someone to be their own brother. Similarly, the order of the arguments of `have_children_together/2` shouldn't matter.

Exercise [ancestor]

It is possible (and often useful) to write recursive rules: rules which reduce the deduction to a base case.

Consider the following definition:

Someone is an ancestor of someone else if they are their direct parent, or if they are the direct parent of one of their ancestors.

Write a rule defining the `ancestor/2` predicate according to the definition.

Beware that the base case of the recursion must be specified first, or the logic engine may enter an endless loop (we will cover the reasons for this next week).

Anonymous Variables

The underscore character can be used instead of variables when we are not interested in their value. This underscore behaves like a regular variable, each of its uses will be treated as a different variable, so the following code does not have the intended effect:

```
grandparent (X, Z) :- parent (X, _), parent (_, Z).
```

Under the hood, each use of `_` will be bound to a distinct fresh variable, using names such as `_1`, `_2`, ...

Exercise [underscore]

Use the underscore wildcard to:

- determine whether Laura has at least one child.
 - add a `human/1` predicate so that *everyone* is a human.
 - define a `is_sibling/1` predicate that checks whether a person has at least one sibling.
-

Exercise [merge]

Consider the following query:

```
findall(Parent, parent(Parent, Child), List).
```

Define a new rule that will help you formulate a query giving the same result, but using `bagof/3` instead, and without using the `^` predicate.

Pattern Matching

Complex terms may appear as parameters to predicates, both in facts and rules:


```
married(couple(george, cecilia)).  
heterosexual(couple(X, Y)) :- male(X), female(Y) ; female(X) , male(Y).
```

as well as in queries:

```
?- married(couple(george, X))  
X = cecilia.  
  
?- married(X)  
X = couple(george, cecilia).
```

When a complex term appears as parameter to a predicate, the part before the parentheses is called *functor* (this does not have the same meaning as in some functional programming languages). In the examples above, `couple` is a functor.

When they appear on the left-hand side of a rule (or in a rule without turnstile), complex parameters behave like an additional unification constraints, so our `heterosexual/1` rule above is in fact equivalent to:

```
heterosexual(Couple) :- Couple = couple(X, Y) , (male(X), female(Y) ; female(X), male(Y)).
```

Pattern matching enables us to do some interesting things:

```
swap(couple(X, Y), couple(Y, X)).
```

```
?- swap(couple(george, cecilia), R).  
R = couple(cecilia, george).
```

As we will soon see, pattern matching is particularly handy when dealing with lists.

Lists

Prolog has a builtin list data structure. They are denoted by square brackets:

```
[a,b,c,d] .
```

Lists are composed of a head and a tail, which is itself a list. To match the head and the tail of the list, you can use the bar operator: `[Head|Tail]` .

The empty list `[]` forms the base case, a list `[a, b, c]` can also be written as `[a|[b|[c|[]]]]`, or even as `[a, b | [c]]` for instance.

Below is an example of a predicate operating over a list.

It binds `R` to the last item in the list, and fails if the list is empty.

```
last1([H], H).  
last1(_|T, R) :- last1(T, R).
```

Exercise `[lists]`

Write Prolog clauses to:

- determine whether a value belongs to a list: `member1(Element, List)`.
- append two lists together `append1(List1, List2, Result)`.
- delete all occurrences of a value in a list: `delete1(Element, List, Result)`.
- rotate all elements of a list to the left (so that the first element of the list will become its last element): `rotate_left(List, Result)`.
- rotate all elements of a list to the right (so that the last element of the list will become its first element): `rotate_right(List, Result)`.

1) Note: the `_1` suffix is to avoid clashing with pre-defined prolog predicates.

2) Note: the symbol for non-unification is `\=`. You can also use the `\+` negation predicate instead: `\+ X = Y`.

3) Hint: you're allowed to define helper predicates! This is particularly useful for `rotate_right/2`.

Exercise `[guess]`

Consider the following clauses:

```
h([], []).  
h([H|T], R) :- h(T, R1), append(R1, [H], R).
```

1. Try to guess the purpose of these clauses.
2. Try to execute `h([1,2,3], R)` by hand.

3. What is the complexity of this predicate? Can you formulate a better version of it with lower complexity?

Hint: use an accumulator.

Bonus Exercises

Write rules to:

1. give all permutations of a given list: `permutations(List, Permutations)` .
2. tell whether a list is a sublist of another: `sublist(Sublist, List)` .

Hint: start by defining a predicate `permutation1/2` which checks if two lists are a permutation of one another.

Don't hesitate to ask for more hints if you need them.